

Credit Card Fraud Detection

Project for Fraud Detection Course

Ricardo Araújo Amorim
up202107843

December 12, 2024

Project Overview

- ▶ **Objective:** Detect fraudulent credit card transactions using machine learning.
- ▶ **Steps in the Process:**
 - ▶ Data Understanding.
 - ▶ Data Preparation.
 - ▶ Clustering.
 - ▶ Modeling.
 - ▶ Evaluation and Results.
- ▶ **Challenges:**
 - ▶ Highly imbalanced dataset (fraud cases = 1.9%).
 - ▶ Complex interactions between features.

Data Understanding: Overview

- ▶ Some attributes were converted to the object type for better handling during data preparation and modeling.
- ▶ **Changed Attributes:**
 - ▶ index: Changed to object.
 - ▶ cc_num: Changed to object.
 - ▶ is_fraud: Changed to object.
 - ▶ zip: Changed to object.
 - ▶ merchant_id: Changed to object.

nangetindex: 30000 entries, 0 to 29999

Data columns (total 25 columns):

#	Column	Non-Null Count	Dtype
0	index	30000 non-null	int64
1	trans_date_trans_time	29900 non-null	object
2	cc_num	30000 non-null	int64
3	device_os	12036 non-null	object
4	merchant	30000 non-null	object
5	amt	29900 non-null	float64
6	trans_num	30000 non-null	object
7	unix_time	30000 non-null	int64
8	is_fraud	30000 non-null	int64
9	first	29990 non-null	object
10	last	29990 non-null	object
11	gender	29990 non-null	object
12	street	29990 non-null	object
13	city	29990 non-null	object
14	zip	29784 non-null	float64
15	job	29784 non-null	object
16	dob	29990 non-null	object
17	category	29401 non-null	object
18	merch_lat	29401 non-null	float64
19	merch_long	29990 non-null	float64
20	merchant_id	29990 non-null	float64
21	lat	10020 non-null	float64
22	long	10020 non-null	float64
23	city_pop	10020 non-null	float64
24	state	10020 non-null	object

Figure: Original Dataset

Correlation Matrix

Key Insights:

- ▶ Analyzed numerical features for linear relationships.
- ▶ Most features show weak or no correlation.
- ▶ Significant correlations observed:
 - ▶ `unix_time` and `index`.
 - ▶ `city_pop` with `lat` and `long`.




Figure: Correlation Matrix

Chi-Square Test Results

Key Insights:

- ▶ Tested independence among categorical variables.
- ▶ Significant dependencies found:
 - ▶ gender and dob.
 - ▶ job and merchant.
- ▶ Highlighted relationships guided feature engineering.
- ▶ Created new interaction variables (e.g., job_age_group).




Figure: Chi-Square Test Results

Data Visualization




Figure: Fraudulent vs Legitimate Transactions




Figure: Transaction Amount Distribution

Data Visualization




Figure: Job Distribution by Age Group




Figure: Distribution of Categories by City

Data Visualization

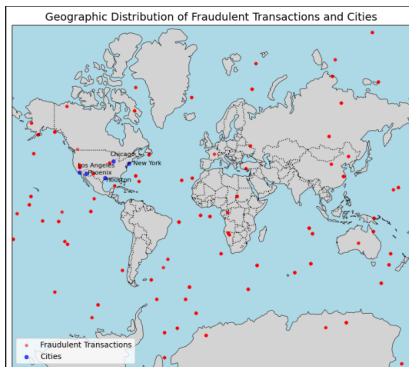


Figure: Geographic Distribution of Fraudulent Transactions




Figure: Transactions by Device OS and Fraud Label

Data Preparation

▶ Dataset Split:

- ▶ Divided into 80% training and 20% testing datasets.
- ▶ Used **stratification** to ensure the same proportion of fraud and non-fraud transactions in both sets.
- ▶ Maintains the balance of the target variable (`is_fraud`) for better model performance and evaluation.

▶ Handling Duplicates:

- ▶ Transaction number (`trans_num`) must be unique to maintain data integrity.
- ▶ Identified duplicate transaction numbers and retained only the record with fewer missing values (NAs).

Feature Engineering

▶ **Age and Age Group Creation:**

- ▶ Calculated age by subtracting the date of birth (`dob`) from the transaction date (`unix_time`).
- ▶ Grouped age into categorical `age_group`:
 - ▶ Bins: *16–30, 31–45, 46–60, 60+*.

▶ **Interaction Features:**

- ▶ Combined `age_group` with `job` to create interaction variables (`job_age_group`).

▶ **Cyclical Temporal Features:**

- ▶ Derived `hour_sin`, `hour_cos` for the hour of the day.

Benefits:

- ▶ Encodes the proximity of hours (e.g., 23:00 and 00:00 are close).
- ▶ Improves model performance by providing meaningful temporal patterns.

Dropped Attributes

- ▶ During data preparation, several attributes were removed to improve model efficiency and focus on relevant features.
- ▶ **Reasons for Removal:**
 - ▶ **Redundancy:**
 - ▶ `index`, `trans_num`: Added no predictive value, purely identifiers.
 - ▶ `trans_date_trans_time`, `transaction_date`: Replaced by derived features like `hour_sin`, `hour_cos`
 - ▶ **Privacy Concerns:**
 - ▶ `first`, `last`, `street`, `state`, `dob`, `cc_num`: Contained personal or sensitive information.

Dropped Attributes

- ▶ **Low Predictive Value:**

- ▶ lat, long, merch_lat, merch_long: Geographic data did not significantly correlate with fraud.
- ▶ zip, merchant, merchant_id: Low variance or redundancy with other features.

- ▶ **Replaced by Derived Features:**

- ▶ age: Replaced by age_group.

Handling Missing Values

- ▶ **Training Data:**

- ▶ **Numeric Variables:**

- ▶ Missing values imputed using the **k-Nearest Neighbors (kNN)** algorithm:
 - ▶ Numeric columns standardized before applying kNN.
 - ▶ Post-imputation, values re-scaled to their original distributions using stored mean and standard deviation.

- ▶ **Categorical Variables:**

- ▶ Handled automatically by kNN, using the **mode of the nearest neighbors**.
 - ▶ Final dataset saved as `X_train_without_missing_values.csv`.

Handling Missing Values

- ▶ **Testing Data:**

- ▶ **Numeric Variables:**

- ▶ Imputed using the **mean values** computed from the training dataset.

- ▶ **Categorical Variables:**

- ▶ Imputed using the **mode values** from the training dataset.

Ensured consistency by aligning imputed values with the training dataset.

One-Hot Encoding

► Purpose:

- Convert categorical variables into a numerical format suitable for machine learning models.

► Process:

- Identified categorical columns using `select_dtypes`.
- Applied **one-hot encoding** to transform these columns:
 - Used `drop_first` for binary categories to avoid multicollinearity.
 - Retained all categories for non-binary columns.
- Ensured the same encoding scheme was applied to both training and testing datasets.

► Output:

- Each categorical column was replaced with multiple binary columns representing the categories.

Normalization and Scaling

► Purpose:

- Normalize numerical variables to ensure all features have comparable scales.
- Prevent variables with larger ranges from dominating the model.

► Why MinMaxScaler?

- Chosen because the distribution of numerical variables is not uniform.
- Scales values to a specified range, typically $[0, 1]$, preserving the original shape of the data distribution.




Figure: Distribution of amt attribute

Normalization and Scaling

► **Benefits of Scaling:**

- Improves the performance of distance-based algorithms (e.g., k-Nearest Neighbors, clustering).
- Ensures gradient-based optimization algorithms converge faster and more reliably.
- Helps prevent bias in models sensitive to variable magnitude.

Clustering Overview

- ▶ **Objective:**

- ▶ Group transactions based on their similarity using unsupervised learning.

- ▶ **Features Used for Clustering:**

- ▶ amt: Transaction amount.
 - ▶ hour_sin, hour_cos: Temporal features representing the hour of the transaction.
 - ▶ city_pop: Population size of the city where the transaction occurred.

- ▶ **Clustering Techniques Explored:**

- ▶ **DBSCAN (Density-Based Spatial Clustering of Applications with Noise).**
 - ▶ **K-Means Clustering.**

DBSCAN Clustering

- ▶ **Approach:**
 - ▶ Identifies clusters based on density of points.
 - ▶ Parameters used:
 - ▶ `eps`: 0.5 (maximum distance between points in a cluster).
 - ▶ `min_samples`: 2 (minimum points to form a dense region).




Figure: DBSCAN Clustering Results

K-Means Clustering

► Approach:

- Groups data into k clusters by minimizing the within-cluster variance.
- Tested different values of k (2 to 10) to find the best clustering structure.




Figure: K-Means Clustering Results (Best k)

Handling Class Imbalance

▶ Class Imbalance:

- ▶ The dataset exhibits significant class imbalance:
 - ▶ Majority class: `is_fraud = 0`.
 - ▶ Minority class: `is_fraud = 1`.
- ▶ Imbalance can lead to biased models with poor recall for fraud detection.

▶ Solution:

- ▶ Combined **SMOTE (Synthetic Minority Oversampling)** and **RandomUnderSampling**.
- ▶ SMOTE increases minority class representation by generating synthetic samples.
- ▶ RandomUnderSampling reduces majority class size, balancing the dataset and preventing computational overhead.

Why SMOTE + RandomUnderSampling?

▶ **Comparison with SMOTE-Tomek:**

- ▶ SMOTE-Tomek removes Tomek links (overlapping samples between classes).
- ▶ Analysis showed no significant class overlap, making SMOTE-Tomek less relevant.

▶ **Advantages of SMOTE + RandomUnderSampling:**

- ▶ Simpler and faster than SMOTE-Tomek.
- ▶ Balances the dataset without unnecessary removal of data points.

▶ **Oversampling and Undersampling Rates:**

- ▶ Oversampling rates tested: 0.5, 0.7 (fractions of the majority class).
- ▶ Undersampling rates tested: 0.8, 0.9 (fractions of the total dataset for the majority class).

Model Pipeline

► Pipeline Steps:

1. Apply **SMOTE** for oversampling the minority class.
2. Apply **RandomUnderSampling** to balance the dataset.
3. Train the model on the balanced dataset.

► Evaluation Metrics:

- **Precision:** Accuracy of fraud predictions.
- **Recall:** Ability to detect fraudulent transactions.
- **F1-Score:** Balance between precision and recall.
- **AUC-ROC:** Overall performance across classification thresholds.

Hyperparameter Tuning: Random Search

- ▶ **Objective:** Improve model performance by finding the best combination of hyperparameters.
- ▶ **Why Random Search?**
 - ▶ More efficient than Grid Search for large hyperparameter spaces.
 - ▶ Allows exploring a wide range of combinations with fewer iterations.
- ▶ **Implementation:**
 - ▶ Used `RandomizedSearchCV` with 5-fold cross-validation.
 - ▶ **What is 5-fold cross-validation?** A technique to evaluate model performance by splitting the data into 5 equally sized subsets, or "folds."
 - ▶ Evaluated models based on the **AUC-ROC** score to handle class imbalance effectively.

Models Overview

► Random Forest:

- Ensemble-based model combining multiple decision trees.
- Effective for imbalanced datasets and interpretable results.

```
Testing Oversampling=0.5, Undersampling=0.8
Best Parameters for this iteration: {'model__n_estimators': 100, 'model__max_depth': 10}
Classification Report:
Confusion Matrix:
[[5886   0]
 [ 114   0]]

Testing Oversampling=0.5, Undersampling=1.0
Best Parameters for this iteration: {'model__n_estimators': 100, 'model__max_depth': 10}
Classification Report:
Confusion Matrix:
[[5886   0]
 [ 114   0]]

Testing Oversampling=0.7, Undersampling=0.8
Best Parameters for this iteration: {'model__n_estimators': 100, 'model__max_depth': 10}
Classification Report:
Confusion Matrix:
[[5885   1]
 [ 114   0]]

Testing Oversampling=0.7, Undersampling=1.0
Best Parameters for this iteration: ('model__n_estimators': 100, 'model__max_depth': 10)
Classification Report:
...
Oversampling Rate: 0.5
Undersampling Rate: 0.8
Best AUC: 0.4489
Submission file created: 'submission_random_search_random_forest.csv'
```

Figure: Random Forest Results Overview

Models Overview

XGBoost:

- ▶ Gradient boosting framework optimized for speed and performance.
- ▶ Excellent at capturing complex, non-linear patterns in data.

```
Testing Oversampling=0.5, Undersampling=0.8
/home/ricardo/Desktop/uni/mestrado/fraude/project/venv/lib/python3.7/site-packages/xgboost/
warnings.warn(
Best Parameters for this iteration: {'model__subsample': 0.6, '
Classification Report:
Confusion Matrix:
[[5150  736]
 [ 105   9]]

Testing Oversampling=0.5, Undersampling=1.0
Best Parameters for this iteration: {'model__subsample': 0.6, '
Classification Report:
Confusion Matrix:
[[5133  753]
 [ 105   9]]

Testing Oversampling=0.7, Undersampling=0.8
Best Parameters for this iteration: {'model__subsample': 0.6, '
Classification Report:
Confusion Matrix:
[[5201  685]
 [ 105   9]]

Testing Oversampling=0.7, Undersampling=1.0
Best Parameters for this iteration: {'model__subsample': 0.6, '
Classification Report:
Confusion Matrix:
[[4990  896]
 [ 105   9]]

...
Oversampling Rate: 0.5
Undersampling Rate: 1.0
Best AUC: 0.4266
Submission file created: 'submission_random_search_xgboost.csv'
```

Figure: XGBoost Results Overview

Models Overview

Decision Tree (Best AUC-ROC Score):

- ▶ Simple and interpretable tree-based model.
- ▶ Tends to overfit but works well with proper pruning and parameter tuning.

```
Testing Oversampling=0.5, Undersampling=0.8
Best Parameters for this iteration: {'model__min_samples_split': 10,
Classification Report:
Confusion Matrix:
[[5282  604]
 [ 97  17]]

Testing Oversampling=0.5, Undersampling=1.0
Best Parameters for this iteration: {'model__min_samples_split': 20,
Classification Report:
Confusion Matrix:
[[3691 2195]
 [ 60  54]]

Testing Oversampling=0.7, Undersampling=0.8
Best Parameters for this iteration: {'model__min_samples_split': 20,
Classification Report:
Confusion Matrix:
[[4537 1349]
 [ 85  29]]

Testing Oversampling=0.7, Undersampling=1.0
Best Parameters for this iteration: {'model__min_samples_split': 10,
Classification Report:
...
Oversampling Rate: 0.5
Undersampling Rate: 1.0
Best AUC: 0.5333
Submission file created: 'submission_decision_tree.csv'
```

Figure: Decision Tree Results Overview

Models Overview

► Multi-Layer Perceptron (MLP):

- Neural network model with hidden layers.
- Effective for capturing non-linear relationships in data.

```
Testing Oversampling=0.5, Undersampling=0.8
Best Parameters for this iteration: {'model_solver': 'sgd'},
Classification Report:
Confusion Matrix:
[[5886   0]
 [ 114   0]]

Testing Oversampling=0.5, Undersampling=1.0
Best Parameters for this iteration: {'model_solver': 'adam'}
Classification Report:
Confusion Matrix:
[[5886   0]
 [ 114   0]]

Testing Oversampling=0.7, Undersampling=0.8
Best Parameters for this iteration: {'model_solver': 'sgd'},
Classification Report:
Confusion Matrix:
[[5886   0]
 [ 114   0]]

Testing Oversampling=0.7, Undersampling=1.0
Best Parameters for this iteration: {'model_solver': 'sgd'},
Classification Report:
...
Oversampling Rate: 0.5
Undersampling Rate: 0.8
Best AUC: 0.5000
Submission file created: 'submission_random_search_mlp.csv'
```

Figure: MLP Results Overview

Models Overview

Support Vector Machine (SVM):

- ▶ Separates data using hyperplanes in high-dimensional space.
- ▶ Effective for smaller datasets and well-separated classes.

```
Testing Oversampling=0.5, Undersampling=0.8
Classification Report:
Confusion Matrix:
[[ 0 5886]
 [ 0 114]]

Testing Oversampling=0.5, Undersampling=1.0
Classification Report:
Confusion Matrix:
[[ 0 5886]
 [ 0 114]]

Testing Oversampling=0.7, Undersampling=0.8
Classification Report:
Confusion Matrix:
[[ 0 5886]
 [ 0 114]]

Testing Oversampling=0.7, Undersampling=1.0
Classification Report:
Confusion Matrix:
[[ 0 5886]
 [ 0 114]]
...
Oversampling Rate: 0.5
Undersampling Rate: 0.8
Best AUC: 0.5000
Submission file created: 'submission_svm.csv'
```

Conclusion

- ▶ The **Decision Tree** performed the best, but with modest results.
- ▶ Valuable insights were gained by following systematic steps.
- ▶ The approach shows potential for better outcomes with real-world data.

Thank you for your attention!