

task2

December 10, 2024

1 Credit Card Fraud Detection

2 Task 2: Predictive Modelling

2.1 Required libraries

```
[97]: import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import numpy as np
import cartopy.crs as ccrs
import cartopy.feature as cfeature
from adjustText import adjust_text
from geopy.distance import geodesic
from sklearn.preprocessing import StandardScaler
from imblearn.under_sampling import RandomUnderSampler
from imblearn.pipeline import Pipeline
from collections import Counter
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.metrics import f1_score
from sklearn.metrics import confusion_matrix, classification_report
from xgboost import XGBClassifier
import pickle
import os
from sklearn.impute import KNNImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.tree import DecisionTreeClassifier
```

2.2

2.3 Load training and test dataset

```
[ ]: with open("variables/X_train.pkl", "rb") as f:
    X_train = pickle.load(f)

with open("variables/y_train.pkl", "rb") as f:
    y_train = pickle.load(f)

with open("variables/X_test.pkl", "rb") as f:
    X_test = pickle.load(f)

with open("variables/y_test.pkl", "rb") as f:
```

```

y_test = pickle.load(f)

with open("variables/kaggle_data.pkl", "rb") as f:
    kaggle_data = pickle.load(f)

with open("variables/index_mapping.pkl", "rb") as f:
    index_mapping = pickle.load(f)

```

```

-----
TypeError                                Traceback (most recent call last)
Cell In[117], line 19
     16 with open("variables/index_mapping.pkl", "rb") as f:
     17     index_mapping = pickle.load(f)
--> 19 y_train.dtype()

TypeError: 'CategoricalDtype' object is not callable

```

2.4 Class Imbalance - SMOTE and model pipeline

To address the significant class imbalance in the dataset, several methods were considered, including **basic SMOTE**, **SMOTE-Tomek**, and **SMOTE combined with RandomUnderSampling**. After evaluating the characteristics of the data, the combination of **SMOTE + RandomUnderSampling** was chosen as the most appropriate technique.

The decision to use **SMOTE + RandomUnderSampling** instead of **SMOTE-Tomek** or **basic SMOTE** was based on the analysis of the dataset's characteristics and the goals of the project.

Class Imbalance The dataset exhibits a significant class imbalance, with very few samples belonging to the minority class (`is_fraud = 1`). Addressing this imbalance is critical to ensure that the model does not become biased towards the majority class (`is_fraud = 0`). SMOTE is effective in resolving this issue by generating synthetic samples for the minority class, thereby improving class representation and model fairness.

Efficiency **SMOTE + RandomUnderSampling** is computationally simpler and faster compared to **SMOTE-Tomek**, as it does not involve the additional step of identifying and removing Tomek links. This makes it a practical choice for the dataset, given the observed lack of significant class overlap.

```

[ ]: oversampling_rates = [0.5,0.6,0.7] # Fraction of majority class
undersampling_rates = [0.8,0.9] # Fraction of total dataset for the majority
    ↪ class

def model_pipeline(model, submission_file, use_random_search=False,
    ↪ param_grid=None, n_iter=10):
    results = []

```

```

best_model = None
best_auc = 0
best_config = None

for oversampling_rate in oversampling_rates:
    for undersampling_rate in undersampling_rates:
        print(f"\nTesting Oversampling={oversampling_rate},\n
↳Undersampling={undersampling_rate}")

        # Define SMOTE and undersampler
        smote = SMOTE(sampling_strategy=oversampling_rate, random_state=42)
        undersampler = \
↳RandomUnderSampler(sampling_strategy=undersampling_rate, random_state=42)

        # Create pipeline
        pipeline = Pipeline(steps=[
            ('smote', smote),
            ('undersampler', undersampler),
            ('model', model)
        ])

        if use_random_search:
            search = RandomizedSearchCV(
                estimator=pipeline,
                param_distributions=param_grid, # Prefix for model params
                n_iter=n_iter,
                scoring='roc_auc',
                cv=5,
                n_jobs=-1,
                random_state=42
            )
            # Train with hyperparameter tuning
            search.fit(X_train, y_train)
            best_pipeline = search.best_estimator_
            best_params = search.best_params_
            print(f"Best Parameters for this iteration: {best_params}")
        else:
            # Train pipeline without hyperparameter search
            pipeline.fit(X_train, y_train)
            best_pipeline = pipeline
            best_params = "Default parameters"

        # Predict on the test set
        y_pred = best_pipeline.predict(X_test)

```

```

        y_pred_proba = best_pipeline.predict_proba(X_test)[: , 1]  # Get
        ↪probabilities for AUC

        # Evaluate model
        print("Classification Report:")
        report = classification_report(y_test, y_pred, output_dict=True,
        ↪zero_division=0)

        # Confusion matrix
        print("Confusion Matrix:")
        print(confusion_matrix(y_test, y_pred))

        # Calculate AUC
        auc_score = roc_auc_score(y_test, y_pred_proba)

        # Store results
        results.append({
            'oversampling_rate': oversampling_rate,
            'undersampling_rate': undersampling_rate,
            'precision': report['1']['precision'],
            'recall': report['1']['recall'],
            'f1_score': report['1']['f1-score'],
            'auc': auc_score
        })

        # Check if current model is the best
        if auc_score > best_auc:
            best_auc = auc_score
            best_model = best_pipeline
            best_config = {
                'oversampling_rate': oversampling_rate,
                'undersampling_rate': undersampling_rate
            }

        # Print best configuration
        print("\nBest Configuration:")
        print(f"Oversampling Rate: {best_config['oversampling_rate']}")
        print(f"Undersampling Rate: {best_config['undersampling_rate']}")
        print(f"Best AUC: {best_auc:.4f}")

        # Predict probabilities for Kaggle submission
        test_probs = best_model.predict_proba(kaggle_data)[: , 1]  # Probabilities
        ↪for class 1 (fraud)

```

```

# Create submission DataFrame
submission = pd.DataFrame({
    'index': index_mapping,
    'is_fraud': test_probs
})

# Save to CSV
submission_file_name = f"{submission_file}.csv"
submission.to_csv(f"submission/{submission_file_name}", index=False)
print(f"Submission file created: '{submission_file_name}'")

return results

```

2.5 Random Forest Classifier

```

[100]: # Train a Random Forest Classifier
rf = RandomForestClassifier(random_state=42)

model_pipeline(rf, "submission_random_forest")

```

Testing Oversampling=0.5, Undersampling=0.8

Classification Report:

Confusion Matrix:

```

[[5882   4]
 [ 114   0]]

```

Testing Oversampling=0.5, Undersampling=1.0

Classification Report:

Confusion Matrix:

```

[[5880   6]
 [ 114   0]]

```

Testing Oversampling=0.7, Undersampling=0.8

Classification Report:

Confusion Matrix:

```

[[5885   1]
 [ 114   0]]

```

Testing Oversampling=0.7, Undersampling=1.0

Classification Report:

Confusion Matrix:

```

[[5882   4]
 [ 114   0]]

```

Best Configuration:

Oversampling Rate: 0.5

Undersampling Rate: 1.0

Best AUC: 0.4655

Submission file created: 'submission_random_forest.csv'

```
[100]: [{'oversampling_rate': 0.5,
        'undersampling_rate': 0.8,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.43127462727494914)},
        {'oversampling_rate': 0.5,
        'undersampling_rate': 1.0,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.4655434244803309)},
        {'oversampling_rate': 0.7,
        'undersampling_rate': 0.8,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.43850781813521233)},
        {'oversampling_rate': 0.7,
        'undersampling_rate': 1.0,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.43363005287598877)}]
```

2.6 Random Search - Random Forest Classifier

```
[105]: rf = RandomForestClassifier(random_state=42)

param_distributions = {
    'model__n_estimators': [100, 200, 500, 1000],
    'model__max_depth': [None, 10, 20, 30],
    'model__min_samples_split': [2, 5, 10],
    'model__min_samples_leaf': [1, 2, 4, 5],
}

model_pipeline(rf, "submission_random_search_random_forest", True, param_distributions)
```

Testing Oversampling=0.5, Undersampling=0.8

Best Parameters for this iteration: {'model__n_estimators': 100, 'model__min_samples_split': 5, 'model__min_samples_leaf': 4, 'model__max_depth': 10}

Classification Report:

Confusion Matrix:

```
[[5886    0]
 [ 114    0]]
```

Testing Oversampling=0.5, Undersampling=1.0

Best Parameters for this iteration: {'model__n_estimators': 100,
'model__min_samples_split': 5, 'model__min_samples_leaf': 4, 'model__max_depth':
10}

Classification Report:

Confusion Matrix:

```
[[5886    0]
 [ 114    0]]
```

Testing Oversampling=0.7, Undersampling=0.8

Best Parameters for this iteration: {'model__n_estimators': 100,
'model__min_samples_split': 5, 'model__min_samples_leaf': 4, 'model__max_depth':
10}

Classification Report:

Confusion Matrix:

```
[[5886    0]
 [ 114    0]]
```

Testing Oversampling=0.7, Undersampling=1.0

Best Parameters for this iteration: {'model__n_estimators': 100,
'model__min_samples_split': 5, 'model__min_samples_leaf': 4, 'model__max_depth':
10}

Classification Report:

Confusion Matrix:

```
[[5886    0]
 [ 114    0]]
```

Best Configuration:

Oversampling Rate: 0.5

Undersampling Rate: 1.0

Best AUC: 0.4935

Submission file created: 'submission_random_search_random_forest.csv'

```
[105]: [{'oversampling_rate': 0.5,
        'undersampling_rate': 0.8,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.47421624908346294)},
        {'oversampling_rate': 0.5,
        'undersampling_rate': 1.0,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
```



```

    'auc': np.float64(0.4935238836132124)}},
{'oversampling_rate': 0.7,
 'undersampling_rate': 0.8,
 'precision': 0.0,
 'recall': 0.0,
 'f1_score': 0.0,
 'auc': np.float64(0.47737271312838675)}},
{'oversampling_rate': 0.7,
 'undersampling_rate': 1.0,
 'precision': 0.0,
 'recall': 0.0,
 'f1_score': 0.0,
 'auc': np.float64(0.4841342227468093)}]}

```

Score on Kaggle: 0.52296

2.7 XGBOOST

```

[106]: xgb = XGBClassifier(n_estimators=500, max_depth=5, learning_rate=0.1,
    ↪random_state=42)

model_pipeline(xgb, "submission_xgboost")

```

Testing Oversampling=0.5, Undersampling=0.8

Classification Report:

Confusion Matrix:

```

[[5841  45]
 [ 114   0]]

```

Testing Oversampling=0.5, Undersampling=1.0

Classification Report:

Confusion Matrix:

```

[[5840  46]
 [ 114   0]]

```

Testing Oversampling=0.7, Undersampling=0.8

Classification Report:

Confusion Matrix:

```

[[5864  22]
 [ 114   0]]

```

Testing Oversampling=0.7, Undersampling=1.0

Classification Report:

Confusion Matrix:

```

[[5851  35]
 [ 114   0]]

```

Best Configuration:
Oversampling Rate: 0.5
Undersampling Rate: 0.8
Best AUC: 0.4529
Submission file created: 'submission_xgboost.csv'

```
[106]: [{ 'oversampling_rate': 0.5,
          'undersampling_rate': 0.8,
          'precision': 0.0,
          'recall': 0.0,
          'f1_score': 0.0,
          'auc': np.float64(0.452863172201656)}},
        { 'oversampling_rate': 0.5,
          'undersampling_rate': 1.0,
          'precision': 0.0,
          'recall': 0.0,
          'f1_score': 0.0,
          'auc': np.float64(0.43907264338215557)}},
        { 'oversampling_rate': 0.7,
          'undersampling_rate': 0.8,
          'precision': 0.0,
          'recall': 0.0,
          'f1_score': 0.0,
          'auc': np.float64(0.4474265429118157)}},
        { 'oversampling_rate': 0.7,
          'undersampling_rate': 1.0,
          'precision': 0.0,
          'recall': 0.0,
          'f1_score': 0.0,
          'auc': np.float64(0.4478035898444719)}]}
```

Score on Kaggle:

2.8 Random Search - XGBOOST

```
[107]: xgb = XGBClassifier(n_estimators=500, max_depth=5, learning_rate=0.1,
    ↪ random_state=42)

param_distributions = {
    'model__n_estimators': [100, 200, 300, 500, 1000],
    'model__learning_rate': [0.01, 0.05, 0.1, 0.2, 0.3],
    'model__max_depth': [3, 5, 7, 10],
    'model__subsample': [0.6, 0.8, 1.0],
    'model__colsample_bytree': [0.6, 0.8, 1.0],
    'model__reg_alpha': [0, 0.1, 1, 10],
    'model__reg_lambda': [1, 10, 50],
    'model__min_child_weight': [1, 3, 5, 7]
```

```
}
```

```
model_pipeline(xgb,"submission_random_search_xgboost",True, param_distributions)
```

Testing Oversampling=0.5, Undersampling=0.8

Best Parameters for this iteration: {'model__subsample': 1.0, 'model__reg_lambda': 1, 'model__reg_alpha': 0, 'model__n_estimators': 100, 'model__min_child_weight': 5, 'model__max_depth': 7, 'model__learning_rate': 0.1, 'model__colsample_bytree': 0.8}

Classification Report:

Confusion Matrix:

```
[[5886    0]
 [ 114    0]]
```

Testing Oversampling=0.5, Undersampling=1.0

Best Parameters for this iteration: {'model__subsample': 1.0, 'model__reg_lambda': 1, 'model__reg_alpha': 0, 'model__n_estimators': 100, 'model__min_child_weight': 5, 'model__max_depth': 7, 'model__learning_rate': 0.1, 'model__colsample_bytree': 0.8}

Classification Report:

Confusion Matrix:

```
[[5886    0]
 [ 114    0]]
```

Testing Oversampling=0.7, Undersampling=0.8

Best Parameters for this iteration: {'model__subsample': 1.0, 'model__reg_lambda': 1, 'model__reg_alpha': 0, 'model__n_estimators': 100, 'model__min_child_weight': 5, 'model__max_depth': 7, 'model__learning_rate': 0.1, 'model__colsample_bytree': 0.8}

Classification Report:

Confusion Matrix:

```
[[5886    0]
 [ 114    0]]
```

Testing Oversampling=0.7, Undersampling=1.0

Best Parameters for this iteration: {'model__subsample': 1.0, 'model__reg_lambda': 1, 'model__reg_alpha': 0, 'model__n_estimators': 100, 'model__min_child_weight': 5, 'model__max_depth': 7, 'model__learning_rate': 0.1, 'model__colsample_bytree': 0.8}

Classification Report:

Confusion Matrix:

```
[[5886    0]
 [ 114    0]]
```

Best Configuration:

Oversampling Rate: 0.7

Undersampling Rate: 0.8

Best AUC: 0.4484

Submission file created: 'submission_random_search_xgboost.csv'

```
[107]: [{'oversampling_rate': 0.5,
        'undersampling_rate': 0.8,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.44771715220773656)},
        {'oversampling_rate': 0.5,
        'undersampling_rate': 1.0,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.4407209196964549)},
        {'oversampling_rate': 0.7,
        'undersampling_rate': 0.8,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.44839896632508897)},
        {'oversampling_rate': 0.7,
        'undersampling_rate': 1.0,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.4454444384832281)}]
```

Score on Kaggle: 0.52380

2.9 Decision Tree

```
[115]: dt = DecisionTreeClassifier(random_state = 42)

parameter_grid = {
    'model__max_depth': [5, 10, 20, 30, 40, 50],
    'model__min_samples_split': [2, 5, 10, 20],
    'model__min_samples_leaf': [1, 2, 4, 5],
    'model__max_leaf_nodes': [None, 10, 20, 50, 100],
    'model__max_features': [1, 2, 3, 4, 5, 6, 7, 8,]
}

model_pipeline(dt, "submission_decision_tree", True, parameter_grid)
```

Testing Oversampling=0.5, Undersampling=0.8

Best Parameters for this iteration: {'model__min_samples_split': 20, 'model__min_samples_leaf': 1, 'model__max_leaf_nodes': 100,

'model__max_features': 4, 'model__max_depth': 50}

Classification Report:

Confusion Matrix:

```
[[5118  768]
 [ 102   12]]
```

Testing Oversampling=0.5, Undersampling=1.0

Best Parameters for this iteration: {'model__min_samples_split': 20,

'model__min_samples_leaf': 1, 'model__max_leaf_nodes': 100,

'model__max_features': 4, 'model__max_depth': 50}

Classification Report:

Confusion Matrix:

```
[[4791 1095]
 [ 100   14]]
```

Testing Oversampling=0.7, Undersampling=0.8

Best Parameters for this iteration: {'model__min_samples_split': 10,

'model__min_samples_leaf': 4, 'model__max_leaf_nodes': None,

'model__max_features': 7, 'model__max_depth': 10}

Classification Report:

Confusion Matrix:

```
[[5850   36]
 [ 112    2]]
```

Testing Oversampling=0.7, Undersampling=1.0

Best Parameters for this iteration: {'model__min_samples_split': 10,

'model__min_samples_leaf': 5, 'model__max_leaf_nodes': None,

'model__max_features': 1, 'model__max_depth': 20}

Classification Report:

Confusion Matrix:

```
[[5213  673]
 [  99   15]]
```

Best Configuration:

Oversampling Rate: 0.7

Undersampling Rate: 1.0

Best AUC: 0.5068

Submission file created: 'submission_decision_tree.csv'

```
[115]: [{'oversampling_rate': 0.5,
        'undersampling_rate': 0.8,
        'precision': 0.015384615384615385,
        'recall': 0.10526315789473684,
        'f1_score': 0.026845637583892617,
        'auc': np.float64(0.4946654565397524)},
        {'oversampling_rate': 0.5,
        'undersampling_rate': 1.0,
```

```

'precision': 0.012623985572587917,
'recall': 0.12280701754385964,
'f1_score': 0.022894521668029435,
'auc': np.float64(0.4907832144070677)},
{'oversampling_rate': 0.7,
 'undersampling_rate': 0.8,
 'precision': 0.05263157894736842,
 'recall': 0.017543859649122806,
 'f1_score': 0.02631578947368421,
 'auc': np.float64(0.4855552276886576)},
{'oversampling_rate': 0.7,
 'undersampling_rate': 1.0,
 'precision': 0.02180232558139535,
 'recall': 0.13157894736842105,
 'f1_score': 0.03740648379052369,
 'auc': np.float64(0.5068322990623008)}}]

```

2.10 Neural Networks - Multi-Layer Perceptron (MLP)

```
[109]: mlp = MLPClassifier(hidden_layer_sizes=(100,), max_iter=300, random_state=42)
```

```
model_pipeline(mlp, "submission_mlp")
```

Testing Oversampling=0.5, Undersampling=0.8

Classification Report:

Confusion Matrix:

```
[[ 0 5886]
 [ 0 114]]
```

Testing Oversampling=0.5, Undersampling=1.0

Classification Report:

Confusion Matrix:

```
[[ 0 5886]
 [ 0 114]]
```

Testing Oversampling=0.7, Undersampling=0.8

Classification Report:

Confusion Matrix:

```
[[5886  0]
 [114  0]]
```

Testing Oversampling=0.7, Undersampling=1.0

Classification Report:

Confusion Matrix:

```
[[5886  0]
```

```
[ 114    0]]
```

```
Best Configuration:
Oversampling Rate: 0.5
Undersampling Rate: 0.8
Best AUC: 0.5000
Submission file created: 'submission_mlp.csv'
```

```
[109]: [{ 'oversampling_rate': 0.5,
          'undersampling_rate': 0.8,
          'precision': 0.019,
          'recall': 1.0,
          'f1_score': 0.03729146221786065,
          'auc': np.float64(0.5)},
        { 'oversampling_rate': 0.5,
          'undersampling_rate': 1.0,
          'precision': 0.019,
          'recall': 1.0,
          'f1_score': 0.03729146221786065,
          'auc': np.float64(0.5)},
        { 'oversampling_rate': 0.7,
          'undersampling_rate': 0.8,
          'precision': 0.0,
          'recall': 0.0,
          'f1_score': 0.0,
          'auc': np.float64(0.5)},
        { 'oversampling_rate': 0.7,
          'undersampling_rate': 1.0,
          'precision': 0.0,
          'recall': 0.0,
          'f1_score': 0.0,
          'auc': np.float64(0.5)}]
```

2.11 Random Search - Neural Networks (MLP)

```
[110]: parameter_grid = {
          'model__hidden_layer_sizes': [(50,), (100,), (100, 50), (150, 100)],
          'model__activation': ['relu', 'tanh'],
          'model__solver': ['adam', 'sgd'],
          'model__alpha': [0.0001, 0.001, 0.01],
          'model__learning_rate': ['constant', 'adaptive']
        }

model_pipeline(mlp, "submission_random_search_mlp", True, parameter_grid)
```

```
Testing Oversampling=0.5, Undersampling=0.8
Best Parameters for this iteration: {'model__solver': 'adam',
```

```
'model__learning_rate': 'adaptive', 'model__hidden_layer_sizes': (100, 50),  
'model__alpha': 0.0001, 'model__activation': 'relu'}
```

Classification Report:

Confusion Matrix:

```
[[5886   0]  
 [ 114   0]]
```

Testing Oversampling=0.5, Undersampling=1.0

Best Parameters for this iteration: {'model__solver': 'adam',

'model__learning_rate': 'adaptive', 'model__hidden_layer_sizes': (100, 50),

'model__alpha': 0.0001, 'model__activation': 'relu'}

Classification Report:

Confusion Matrix:

```
[[  0 5886]  
 [  0 114]]
```

Testing Oversampling=0.7, Undersampling=0.8

Best Parameters for this iteration: {'model__solver': 'adam',

'model__learning_rate': 'constant', 'model__hidden_layer_sizes': (50,),

'model__alpha': 0.0001, 'model__activation': 'relu'}

Classification Report:

Confusion Matrix:

```
[[ 365 5521]  
 [  7 107]]
```

Testing Oversampling=0.7, Undersampling=1.0

Best Parameters for this iteration: {'model__solver': 'adam',

'model__learning_rate': 'constant', 'model__hidden_layer_sizes': (50,),

'model__alpha': 0.0001, 'model__activation': 'relu'}

Classification Report:

Confusion Matrix:

```
[[5886   0]  
 [ 114   0]]
```

Best Configuration:

Oversampling Rate: 0.7

Undersampling Rate: 0.8

Best AUC: 0.5003

Submission file created: 'submission_random_search_mlp.csv'

```
[110]: [{'oversampling_rate': 0.5,  
        'undersampling_rate': 0.8,  
        'precision': 0.0,  
        'recall': 0.0,  
        'f1_score': 0.0,  
        'auc': np.float64(0.5)},  
        {'oversampling_rate': 0.5,
```



```

'undersampling_rate': 1.0,
'precision': 0.019,
'recall': 1.0,
'f1_score': 0.03729146221786065,
'auc': np.float64(0.5)}},
{'oversampling_rate': 0.7,
'undersampling_rate': 0.8,
'precision': 0.019012082444918265,
'recall': 0.9385964912280702,
'f1_score': 0.03726924416579589,
'auc': np.float64(0.5003040220326556)}},
{'oversampling_rate': 0.7,
'undersampling_rate': 1.0,
'precision': 0.0,
'recall': 0.0,
'f1_score': 0.0,
'auc': np.float64(0.5)}}]

```

2.12 Support Vector Machine (SVM)

```
[111]: svm = SVC(kernel='rbf', probability=True, random_state=42)
```

```
model_pipeline(svm, "submission_svm")
```

```

-----
NameError                                Traceback (most recent call last)
Cell In[111], line 1
----> 1 aa
      2 svm = SVC(kernel='rbf', probability=True, random_state=42)
      3 #svm.fit(X_train, y_train)

NameError: name 'aa' is not defined

```