

task2

December 10, 2024

1 Credit Card Fraud Detection

2 Task 2: Predictive Modelling

2.1 Required libraries

```
[1]: import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import numpy as np
import cartopy.crs as ccrs
import cartopy.feature as cfeature
from adjustText import adjust_text
from geopy.distance import geodesic
from sklearn.preprocessing import StandardScaler
from imblearn.under_sampling import RandomUnderSampler
from imblearn.pipeline import Pipeline
from collections import Counter
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split
from sklearn.cluster import DBSCAN
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.metrics import f1_score
from sklearn.metrics import confusion_matrix, classification_report
from xgboost import XGBClassifier
import pickle
import os
from sklearn.impute import KNNImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.tree import DecisionTreeClassifier
```

2.2 Load training and test dataset

```
[2]: with open("variables/X_train.pkl", "rb") as f:
    X_train = pickle.load(f)

with open("variables/y_train.pkl", "rb") as f:
    y_train = pickle.load(f)

with open("variables/X_test.pkl", "rb") as f:
    X_test = pickle.load(f)

with open("variables/y_test.pkl", "rb") as f:
    y_test = pickle.load(f)
```

```

with open("variables/kaggle_data.pkl", "rb") as f:
    kaggle_data = pickle.load(f)

with open("variables/index_mapping.pkl", "rb") as f:
    index_mapping = pickle.load(f)

```

2.3 Class Imbalance - SMOTE and model pipeline

To address the significant class imbalance in the dataset, several methods were considered, including **basic SMOTE**, **SMOTE-Tomek**, and **SMOTE combined with RandomUnderSampling**. After evaluating the characteristics of the data, the combination of **SMOTE + RandomUnderSampling** was chosen as the most appropriate technique.

The decision to use **SMOTE + RandomUnderSampling** instead of **SMOTE-Tomek** or **basic SMOTE** was based on the analysis of the dataset's characteristics and the goals of the project.

Class Imbalance The dataset exhibits a significant class imbalance, with very few samples belonging to the minority class (`is_fraud = 1`). Addressing this imbalance is critical to ensure that the model does not become biased towards the majority class (`is_fraud = 0`). SMOTE is effective in resolving this issue by generating synthetic samples for the minority class, thereby improving class representation and model fairness.

Efficiency **SMOTE + RandomUnderSampling** is computationally simpler and faster compared to **SMOTE-Tomek**, as it does not involve the additional step of identifying and removing Tomek links. This makes it a practical choice for the dataset, given the observed lack of significant class overlap.

```

[3]: oversampling_rates = [0.5,0.7] # Fraction of majority class
undersampling_rates = [0.8,1.0] # Fraction of total dataset for the majority
    ↪ class

def model_pipeline(model, submission_file, use_random_search=False,
    ↪ param_grid=None, n_iter=10):
    results = []
    best_model = None
    best_auc = 0
    best_config = None

    for oversampling_rate in oversampling_rates:
        for undersampling_rate in undersampling_rates:
            print(f"\nTesting Oversampling={oversampling_rate},
    ↪ Undersampling={undersampling_rate}")

            # Define SMOTE and undersampler
            smote = SMOTE(sampling_strategy=oversampling_rate, random_state=42)

```

```

undersampler = ␣
↪RandomUnderSampler(sampling_strategy=undersampling_rate, random_state=42)

# Create pipeline
pipeline = Pipeline(steps=[
    ('smote', smote),
    ('undersampler', undersampler),
    ('model', model)
])

if use_random_search:
    search = RandomizedSearchCV(
        estimator=pipeline,
        param_distributions=param_grid, # Prefix for model params
        n_iter=n_iter,
        scoring='roc_auc',
        cv=5,
        n_jobs=-1,
        random_state=42
    )
    # Train with hyperparameter tuning
    search.fit(X_train, y_train)
    best_pipeline = search.best_estimator_
    best_params = search.best_params_
    print(f"Best Parameters for this iteration: {best_params}")
else:
    # Train pipeline without hyperparameter search
    pipeline.fit(X_train, y_train)
    best_pipeline = pipeline
    best_params = "Default parameters"

# Predict on the test set
y_pred = best_pipeline.predict(X_test)
y_pred_proba = best_pipeline.predict_proba(X_test)[:, 1] # Get ␣
↪probabilities for AUC

# Evaluate model
print("Classification Report:")
report = classification_report(y_test, y_pred, output_dict=True, ␣
↪zero_division=0)

# Confusion matrix
print("Confusion Matrix:")

```

```

print(confusion_matrix(y_test, y_pred))

# Calculate AUC
auc_score = roc_auc_score(y_test, y_pred_proba)

# Store results
results.append({
    'oversampling_rate': oversampling_rate,
    'undersampling_rate': undersampling_rate,
    'precision': report['1']['precision'],
    'recall': report['1']['recall'],
    'f1_score': report['1']['f1-score'],
    'auc': auc_score
})

# Check if current model is the best
if auc_score > best_auc:
    best_auc = auc_score
    best_model = best_pipeline
    best_config = {
        'oversampling_rate': oversampling_rate,
        'undersampling_rate': undersampling_rate
    }

# Print best configuration
print("\nBest Configuration:")
print(f"Oversampling Rate: {best_config['oversampling_rate']}")
print(f"Undersampling Rate: {best_config['undersampling_rate']}")
print(f"Best AUC: {best_auc:.4f}")

# Predict probabilities for Kaggle submission
test_probs = best_model.predict_proba(kaggle_data)[: , 1] # Probabilities
↳ for class 1 (fraud)

# Create submission DataFrame
submission = pd.DataFrame({
    'index': index_mapping,
    'is_fraud': test_probs
})

# Save to CSV
submission_file_name = f"{submission_file}.csv"
submission.to_csv(f"submission/{submission_file_name}", index=False)
print(f"Submission file created: '{submission_file_name}'")

return results

```

2.4 Random Forest Classifier

```
[4]: # Train a Random Forest Classifier
rf = RandomForestClassifier(random_state=42)

model_pipeline(rf,"submission_random_forest")
```

Testing Oversampling=0.5, Undersampling=0.8

Classification Report:

Confusion Matrix:

```
[[5883   3]
 [ 114   0]]
```

Testing Oversampling=0.5, Undersampling=1.0

Classification Report:

Confusion Matrix:

```
[[5884   2]
 [ 114   0]]
```

Testing Oversampling=0.7, Undersampling=0.8

Classification Report:

Confusion Matrix:

```
[[5885   1]
 [ 114   0]]
```

Testing Oversampling=0.7, Undersampling=1.0

Classification Report:

Confusion Matrix:

```
[[5886   0]
 [ 114   0]]
```

Best Configuration:

Oversampling Rate: 0.5

Undersampling Rate: 1.0

Best AUC: 0.4390

Submission file created: 'submission_random_forest.csv'

```
[4]: [{'oversampling_rate': 0.5,
      'undersampling_rate': 0.8,
      'precision': 0.0,
      'recall': 0.0,
      'f1_score': 0.0,
      'auc': np.float64(0.4270488700514453)},
      {'oversampling_rate': 0.5,
      'undersampling_rate': 1.0,
      'precision': 0.0,
      'recall': 0.0,
```

```

    'f1_score': 0.0,
    'auc': np.float64(0.4389839702892978)}},
{'oversampling_rate': 0.7,
 'undersampling_rate': 0.8,
 'precision': 0.0,
 'recall': 0.0,
 'f1_score': 0.0,
 'auc': np.float64(0.4185668043707639)}},
{'oversampling_rate': 0.7,
 'undersampling_rate': 1.0,
 'precision': 0.0,
 'recall': 0.0,
 'f1_score': 0.0,
 'auc': np.float64(0.41083212618702725)}}]

```

2.5 Random Search - Random Forest Classifier

```

[10]: rf = RandomForestClassifier(random_state=42)

param_distributions = {
    'model__n_estimators': [100, 200, 500,1000],
    'model__max_depth': [None, 10, 20, 30],
    'model__min_samples_split': [2, 5, 10],
    'model__min_samples_leaf': [1, 2, 4, 5],
}

model_pipeline(rf,"submission_random_search_random_forest",True,param_distributions)

```

Testing Oversampling=0.5, Undersampling=0.8

Best Parameters for this iteration: {'model__n_estimators': 100, 'model__min_samples_split': 5, 'model__min_samples_leaf': 4, 'model__max_depth': 10}

Classification Report:

Confusion Matrix:

```

[[5886   0]
 [ 114   0]]

```

Testing Oversampling=0.5, Undersampling=1.0

Best Parameters for this iteration: {'model__n_estimators': 100, 'model__min_samples_split': 5, 'model__min_samples_leaf': 4, 'model__max_depth': 10}

Classification Report:

Confusion Matrix:

```

[[5886   0]
 [ 114   0]]

```

Testing Oversampling=0.7, Undersampling=0.8

Best Parameters for this iteration: {'model__n_estimators': 100,
'model__min_samples_split': 5, 'model__min_samples_leaf': 4, 'model__max_depth':
10}

Classification Report:

Confusion Matrix:

```
[[5885   1]
 [ 114   0]]
```

Testing Oversampling=0.7, Undersampling=1.0

Best Parameters for this iteration: {'model__n_estimators': 100,
'model__min_samples_split': 5, 'model__min_samples_leaf': 4, 'model__max_depth':
10}

Classification Report:

Confusion Matrix:

```
[[5886   0]
 [ 113   1]]
```

Best Configuration:

Oversampling Rate: 0.5

Undersampling Rate: 0.8

Best AUC: 0.4489

Submission file created: 'submission_random_search_random_forest.csv'

```
[10]: [{'oversampling_rate': 0.5,
        'undersampling_rate': 0.8,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.44885947028631723)},
        {'oversampling_rate': 0.5,
        'undersampling_rate': 1.0,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.4375748877801026)},
        {'oversampling_rate': 0.7,
        'undersampling_rate': 0.8,
        'precision': 0.0,
        'recall': 0.0,
        'f1_score': 0.0,
        'auc': np.float64(0.43664419288111544)},
        {'oversampling_rate': 0.7,
        'undersampling_rate': 1.0,
        'precision': 1.0,
        'recall': 0.008771929824561403,
        'f1_score': 0.017391304347826087,
        'auc': np.float64(0.43613674434131544)}]]
```


2.6 XGBOOST

```
[ ]: xgb = XGBClassifier(n_estimators=500, max_depth=5, learning_rate=0.1,
    ↪random_state=42)

model_pipeline(xgb, "submission_xgboost")
```

Testing Oversampling=0.5, Undersampling=0.8

Classification Report:

Confusion Matrix:

```
[[4226 1660]
 [ 93   21]]
```

Testing Oversampling=0.5, Undersampling=1.0

Classification Report:

Confusion Matrix:

```
[[4215 1671]
 [ 91   23]]
```

Testing Oversampling=0.7, Undersampling=0.8

Classification Report:

Confusion Matrix:

```
[[4350 1536]
 [ 91   23]]
```

Testing Oversampling=0.7, Undersampling=1.0

Classification Report:

Confusion Matrix:

```
[[4307 1579]
 [ 92   22]]
```

Best Configuration:

Oversampling Rate: 0.7

Undersampling Rate: 1.0

Best AUC: 0.4238

Submission file created: 'submission_xgboost.csv'

```
[ ]: [{'oversampling_rate': 0.5,
      'undersampling_rate': 0.8,
      'precision': 0.012492563950029744,
      'recall': 0.18421052631578946,
      'f1_score': 0.0233983286908078,
      'auc': np.float64(0.408907398465583)},
      {'oversampling_rate': 0.5,
      'undersampling_rate': 1.0,
      'precision': 0.01357733175914994,
```

```

'recall': 0.20175438596491227,
'f1_score': 0.025442477876106196,
'auc': np.float64(0.41800719518810625)},
{'oversampling_rate': 0.7,
 'undersampling_rate': 0.8,
 'precision': 0.014753046824887749,
 'recall': 0.20175438596491227,
 'f1_score': 0.02749551703526599,
 'auc': np.float64(0.4136182496676622)},
{'oversampling_rate': 0.7,
 'undersampling_rate': 1.0,
 'precision': 0.013741411617738912,
 'recall': 0.19298245614035087,
 'f1_score': 0.02565597667638484,
 'auc': np.float64(0.4237552980310102)}]

```

2.7 Random Search - XGBOOST

```

[ ]: xgb = XGBClassifier(n_estimators=500, max_depth=5, learning_rate=0.1,
    ↪random_state=42)

param_distributions = {
    'model__n_estimators': [100, 200, 300, 500, 1000],
    'model__learning_rate': [0.01, 0.05, 0.1, 0.2, 0.3],
    'model__max_depth': [3, 5, 7, 10],
    'model__subsample': [0.6, 0.8, 1.0],
    'model__colsample_bytree': [0.6, 0.8, 1.0],
    'model__reg_alpha': [0, 0.1, 1, 10],
    'model__reg_lambda': [1, 10, 50],
    'model__min_child_weight': [1, 3, 5, 7]
}

model_pipeline(xgb, "submission_random_search_xgboost", True, param_distributions)

```

Testing Oversampling=0.5, Undersampling=0.8

/home/ricardo/Desktop/uni/mestrado/fraude/project/venv/lib/python3.12/site-packages/joblib/externals/loky/process_executor.py:752: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.

warnings.warn(

Best Parameters for this iteration: {'model__subsample': 0.6, 'model__reg_lambda': 10, 'model__reg_alpha': 0, 'model__n_estimators': 100, 'model__min_child_weight': 3, 'model__max_depth': 7, 'model__learning_rate': 0.01, 'model__colsample_bytree': 0.8}

Classification Report:

Confusion Matrix:

```
[[5150  736]
 [ 105    9]]
```

Testing Oversampling=0.5, Undersampling=1.0

Best Parameters for this iteration: {'model__subsample': 0.6, 'model__reg_lambda': 10, 'model__reg_alpha': 0, 'model__n_estimators': 100, 'model__min_child_weight': 3, 'model__max_depth': 7, 'model__learning_rate': 0.01, 'model__colsample_bytree': 0.8}

Classification Report:

Confusion Matrix:

```
[[5133  753]
 [ 105    9]]
```

Testing Oversampling=0.7, Undersampling=0.8

Best Parameters for this iteration: {'model__subsample': 0.6, 'model__reg_lambda': 10, 'model__reg_alpha': 0, 'model__n_estimators': 100, 'model__min_child_weight': 3, 'model__max_depth': 7, 'model__learning_rate': 0.01, 'model__colsample_bytree': 0.8}

Classification Report:

Confusion Matrix:

```
[[5201  685]
 [ 105    9]]
```

Testing Oversampling=0.7, Undersampling=1.0

Best Parameters for this iteration: {'model__subsample': 0.6, 'model__reg_lambda': 10, 'model__reg_alpha': 0, 'model__n_estimators': 100, 'model__min_child_weight': 3, 'model__max_depth': 7, 'model__learning_rate': 0.01, 'model__colsample_bytree': 0.8}

Classification Report:

Confusion Matrix:

```
[[4990  896]
 [ 105    9]]
```

Best Configuration:

Oversampling Rate: 0.5

Undersampling Rate: 1.0

Best AUC: 0.4266

Submission file created: 'submission_random_search_xgboost.csv'

```
[ ]: [{'oversampling_rate': 0.5,
      'undersampling_rate': 0.8,
      'precision': 0.012080536912751677,
      'recall': 0.07894736842105263,
      'f1_score': 0.020954598370197905,
      'auc': np.float64(0.4217083653748711)},
      {'oversampling_rate': 0.5,
```

```

'undersampling_rate': 1.0,
'precision': 0.011811023622047244,
'recall': 0.07894736842105263,
'f1_score': 0.02054794520547945,
'auc': np.float64(0.42658613063409456)},
{'oversampling_rate': 0.7,
'undersampling_rate': 0.8,
'precision': 0.012968299711815562,
'recall': 0.07894736842105263,
'f1_score': 0.022277227722772276,
'auc': np.float64(0.4223812376677338)},
{'oversampling_rate': 0.7,
'undersampling_rate': 1.0,
'precision': 0.009944751381215469,
'recall': 0.07894736842105263,
'f1_score': 0.017664376840039256,
'auc': np.float64(0.4227784037054921)}}]

```

2.8 Decision Tree

```

[ ]: dt = DecisionTreeClassifier(random_state = 42)

parameter_grid = {
    'model__max_depth': [5, 10, 20, 30, 40, 50],
    'model__min_samples_split': [2, 5, 10, 20],
    'model__min_samples_leaf': [1, 2, 4, 5],
    'model__max_leaf_nodes': [None, 10, 20, 50, 100],
    'model__max_features': [1, 2, 3, 4, 5, 6, 7, 8,]
}

model_pipeline(dt,"submission_decision_tree",True, parameter_grid)

```

Testing Oversampling=0.5, Undersampling=0.8

Best Parameters for this iteration: {'model__min_samples_split': 10, 'model__min_samples_leaf': 4, 'model__max_leaf_nodes': None, 'model__max_features': 7, 'model__max_depth': 10}

Classification Report:

Confusion Matrix:

```

[[5282  604]
 [  97   17]]

```

Testing Oversampling=0.5, Undersampling=1.0

Best Parameters for this iteration: {'model__min_samples_split': 20, 'model__min_samples_leaf': 2, 'model__max_leaf_nodes': 50, 'model__max_features': 6, 'model__max_depth': 10}

Classification Report:

Confusion Matrix:

```
[[3691 2195]
 [ 60  54]]
```

Testing Oversampling=0.7, Undersampling=0.8

Best Parameters for this iteration: {'model__min_samples_split': 20,
'model__min_samples_leaf': 2, 'model__max_leaf_nodes': 50,
'model__max_features': 6, 'model__max_depth': 10}

Classification Report:

Confusion Matrix:

```
[[4537 1349]
 [ 85  29]]
```

Testing Oversampling=0.7, Undersampling=1.0

Best Parameters for this iteration: {'model__min_samples_split': 10,
'model__min_samples_leaf': 5, 'model__max_leaf_nodes': None,
'model__max_features': 1, 'model__max_depth': 20}

Classification Report:

Confusion Matrix:

```
[[ 803 5083]
 [ 19  95]]
```

Best Configuration:

Oversampling Rate: 0.5

Undersampling Rate: 1.0

Best AUC: 0.5333

Submission file created: 'submission_decision_tree.csv'

```
[ ]: [{'oversampling_rate': 0.5,
      'undersampling_rate': 0.8,
      'precision': 0.027375201288244767,
      'recall': 0.14912280701754385,
      'f1_score': 0.04625850340136054,
      'auc': np.float64(0.5215647000614005)},
      {'oversampling_rate': 0.5,
      'undersampling_rate': 1.0,
      'precision': 0.02401067140951534,
      'recall': 0.47368421052631576,
      'f1_score': 0.045704612780363946,
      'auc': np.float64(0.5332613516461898)},
      {'oversampling_rate': 0.7,
      'undersampling_rate': 0.8,
      'precision': 0.02104499274310595,
      'recall': 0.2543859649122807,
      'f1_score': 0.0388739946380697,
      'auc': np.float64(0.4858979976274359)},
      {'oversampling_rate': 0.7,
      'undersampling_rate': 1.0,
```

```
'precision': 0.018346852066434917,  
'recall': 0.8333333333333334,  
'f1_score': 0.035903250188964474,  
'auc': np.float64(0.4907176410274753)]}]
```

2.9 Neural Networks - Multi-Layer Perceptron (MLP)

```
[9]: mlp = MLPClassifier(max_iter=300, random_state=42)
```

```
model_pipeline(mlp, "submission_mlp")
```

Testing Oversampling=0.5, Undersampling=0.8

Classification Report:

Confusion Matrix:

```
[[5886   0]  
 [ 114   0]]
```

Testing Oversampling=0.5, Undersampling=1.0

/home/ricardo/Desktop/uni/mestrado/fraude/project/venv/lib/python3.12/site-packages/sklearn/neural_network/_multilayer_perceptron.py:690:

ConvergenceWarning: Stochastic Optimizer: Maximum iterations (300) reached and the optimization hasn't converged yet.

warnings.warn(

Classification Report:

Confusion Matrix:

```
[[5886   0]  
 [ 114   0]]
```

Testing Oversampling=0.7, Undersampling=0.8

/home/ricardo/Desktop/uni/mestrado/fraude/project/venv/lib/python3.12/site-packages/sklearn/neural_network/_multilayer_perceptron.py:690:

ConvergenceWarning: Stochastic Optimizer: Maximum iterations (300) reached and the optimization hasn't converged yet.

warnings.warn(

Classification Report:

Confusion Matrix:

```
[[5886   0]  
 [ 114   0]]
```

Testing Oversampling=0.7, Undersampling=1.0

Classification Report:

Confusion Matrix:

```
[[5886   0]
```

```
[ 114    0]]
```

Best Configuration:

Oversampling Rate: 0.5

Undersampling Rate: 0.8

Best AUC: 0.5000

Submission file created: 'submission_mlp.csv'

/home/ricardo/Desktop/uni/mestrado/fraude/project/venv/lib/python3.12/site-packages/sklearn/neural_network/_multilayer_perceptron.py:690:

ConvergenceWarning: Stochastic Optimizer: Maximum iterations (300) reached and the optimization hasn't converged yet.

warnings.warn(

```
[9]: [{'oversampling_rate': 0.5,
      'undersampling_rate': 0.8,
      'precision': 0.0,
      'recall': 0.0,
      'f1_score': 0.0,
      'auc': np.float64(0.5)},
      {'oversampling_rate': 0.5,
      'undersampling_rate': 1.0,
      'precision': 0.0,
      'recall': 0.0,
      'f1_score': 0.0,
      'auc': np.float64(0.5)},
      {'oversampling_rate': 0.7,
      'undersampling_rate': 0.8,
      'precision': 0.0,
      'recall': 0.0,
      'f1_score': 0.0,
      'auc': np.float64(0.5)},
      {'oversampling_rate': 0.7,
      'undersampling_rate': 1.0,
      'precision': 0.0,
      'recall': 0.0,
      'f1_score': 0.0,
      'auc': np.float64(0.5)}]
```

2.10 Random Search - Neural Networks (MLP)

```
[8]: mlp = MLPClassifier( max_iter=300, random_state=42,early_stopping=True)

parameter_grid = {
    'model__hidden_layer_sizes': [(50,), (100,), (100, 50)],
    'model__activation': ['relu', 'tanh'],
    'model__solver': ['adam', 'sgd'],
```

```

        'model__alpha': [0.0001, 0.001],
        'model__learning_rate': ['constant', 'adaptive'],
        'model__max_iter': [300, 500, 1000]
    }

model_pipeline(mlp, "submission_random_search_mlp", True, parameter_grid)

```

Testing Oversampling=0.5, Undersampling=0.8

Best Parameters for this iteration: {'model__solver': 'sgd', 'model__max_iter': 300, 'model__learning_rate': 'constant', 'model__hidden_layer_sizes': (100, 50), 'model__alpha': 0.0001, 'model__activation': 'tanh'}

Classification Report:

Confusion Matrix:

```

[[5886   0]
 [ 114   0]]

```

Testing Oversampling=0.5, Undersampling=1.0

Best Parameters for this iteration: {'model__solver': 'adam', 'model__max_iter': 1000, 'model__learning_rate': 'adaptive', 'model__hidden_layer_sizes': (50,), 'model__alpha': 0.0001, 'model__activation': 'tanh'}

Classification Report:

Confusion Matrix:

```

[[5886   0]
 [ 114   0]]

```

Testing Oversampling=0.7, Undersampling=0.8

Best Parameters for this iteration: {'model__solver': 'sgd', 'model__max_iter': 300, 'model__learning_rate': 'constant', 'model__hidden_layer_sizes': (100, 50), 'model__alpha': 0.0001, 'model__activation': 'tanh'}

Classification Report:

Confusion Matrix:

```

[[5886   0]
 [ 114   0]]

```

Testing Oversampling=0.7, Undersampling=1.0

Best Parameters for this iteration: {'model__solver': 'sgd', 'model__max_iter': 300, 'model__learning_rate': 'constant', 'model__hidden_layer_sizes': (100, 50), 'model__alpha': 0.0001, 'model__activation': 'tanh'}

Classification Report:

Confusion Matrix:

```

[[5886   0]
 [ 114   0]]

```

Best Configuration:

Oversampling Rate: 0.5

Undersampling Rate: 0.8

Best AUC: 0.5000

Submission file created: 'submission_random_search_mlp.csv'

```
[8]: [{'oversampling_rate': 0.5,
      'undersampling_rate': 0.8,
      'precision': 0.0,
      'recall': 0.0,
      'f1_score': 0.0,
      'auc': np.float64(0.5)},
      {'oversampling_rate': 0.5,
      'undersampling_rate': 1.0,
      'precision': 0.0,
      'recall': 0.0,
      'f1_score': 0.0,
      'auc': np.float64(0.5)},
      {'oversampling_rate': 0.7,
      'undersampling_rate': 0.8,
      'precision': 0.0,
      'recall': 0.0,
      'f1_score': 0.0,
      'auc': np.float64(0.5)},
      {'oversampling_rate': 0.7,
      'undersampling_rate': 1.0,
      'precision': 0.0,
      'recall': 0.0,
      'f1_score': 0.0,
      'auc': np.float64(0.5)}]
```

2.11 Support Vector Machine (SVM)

```
[ ]: svm = SVC(kernel='rbf', probability=True, random_state=42)

model_pipeline(svm, "submission_svm")
```

Testing Oversampling=0.5, Undersampling=0.8
Classification Report:
Confusion Matrix:
[[0 5886]
 [0 114]]

Testing Oversampling=0.5, Undersampling=1.0
Classification Report:
Confusion Matrix:
[[0 5886]
 [0 114]]

Testing Oversampling=0.7, Undersampling=0.8

Classification Report:

Confusion Matrix:

```
[[ 0 5886]
 [ 0 114]]
```

Testing Oversampling=0.7, Undersampling=1.0

Classification Report:

Confusion Matrix:

```
[[ 0 5886]
 [ 0 114]]
```

Best Configuration:

Oversampling Rate: 0.5

Undersampling Rate: 0.8

Best AUC: 0.5000

Submission file created: 'submission_svm.csv'

```
[ ]: [{'oversampling_rate': 0.5,
      'undersampling_rate': 0.8,
      'precision': 0.019,
      'recall': 1.0,
      'f1_score': 0.03729146221786065,
      'auc': np.float64(0.5)},
      {'oversampling_rate': 0.5,
      'undersampling_rate': 1.0,
      'precision': 0.019,
      'recall': 1.0,
      'f1_score': 0.03729146221786065,
      'auc': np.float64(0.5)},
      {'oversampling_rate': 0.7,
      'undersampling_rate': 0.8,
      'precision': 0.019,
      'recall': 1.0,
      'f1_score': 0.03729146221786065,
      'auc': np.float64(0.5)},
      {'oversampling_rate': 0.7,
      'undersampling_rate': 1.0,
      'precision': 0.019,
      'recall': 1.0,
      'f1_score': 0.03729146221786065,
      'auc': np.float64(0.5)}]
```

3 Conclusion

After analyzing and experimenting with various models, the **Decision Tree** appeared to perform the best in the current setup, even though the overall score was not particularly high.

I made an effort to follow the appropriate steps required for success if the dataset were real,

and the process provided valuable insights into model development. While the **Decision Tree** outperformed the other tested algorithms, its performance was still relatively modest. **Decision Trees** are known for capturing complex relationships within data, which made them a suitable choice for this experiment. I aimed to adhere to best practices, including data preprocessing, feature selection, and hyperparameter tuning, and applied methods that are scalable to real-world datasets.

Although the final results were not exceptional, this exercise was a meaningful attempt to systematically develop and evaluate models. With a real-world dataset, this approach could potentially lead to more promising outcomes.